

[Up to home page:](#)

[Get 128bit.tgz:](#)

[Documentation for single precision functions.](#)
[Documentation for double precision functions.](#)
[Documentation for 80-bit long double functions.](#)
[Documentation for 128-bit long double functions.](#)
[Documentation for extended precision functions.](#)

128-bit Long Double Precision Functions

Select function name for additional information. For other precisions, see the archives and descriptions listed above.

- [acoshll](#), [Inverse hyperbolic cosine](#)
- [asinhll](#), [Inverse hyperbolic sine](#)
- [asinll](#), [Inverse circular sine](#)
- [acosll](#), [Inverse circular cosine](#)
- [atanhll](#), [Inverse hyperbolic tangent](#)
- [atanll](#), [Inverse circular tangent](#)
- [atan2ll](#), [Quadrant correct inverse circular tangent](#)
- [cbrtll](#), [Cube root](#)
- [coshll](#), [Hyperbolic cosine](#)
- [exp10ll](#), [Base 10 exponential function](#)
- [exp2ll](#), [Base 2 exponential function](#)
- [expll](#), [Exponential function](#)
- [expm1ll](#), [Exponential function, minus 1](#)
- [ceilll](#), [Round up to integer](#)
- [floorll](#), [Round down to integer](#)
- [frexp1ll](#), [Extract exponent and significand](#)
- [ldexp1ll](#), [Apply exponent](#)
- [fabsll](#), [Absolute value](#)
- [signbitll](#), [Extract sign](#)
- [isnanll](#), [Test for not a number](#)
- [isfinitell](#), [Test for infinity](#)
- [ieee](#), [Extended precision arithmetic](#)
- [j0ll](#), [Bessel function, first kind, order 0](#)
- [y0ll](#), [Bessel function, second kind, order 0](#)
- [j1ll](#), [Bessel function, first kind, order 1](#)
- [y1ll](#), [Bessel function, second kind, order 1](#)
- [jnll](#), [Bessel function, first kind, order 1](#)
- [lgammall](#), [Logarithm of gamma function](#)
- [log10ll](#), [Common logarithm](#)
- [log1pll](#), [Relative error logarithm](#)
- [log2ll](#), [Base 2 logarithm](#)
- [logll](#), [Natural logarithm](#)
- [ndtrll](#), [Normal distribution function](#)
- [erfl](#), [Error function](#)
- [erfcfl](#), [Error function](#)
- [mtherr](#), [Error handling](#)
- [polevll](#), [Evaluate polynomial](#)
- [plevll](#), [Evaluate polynomial](#)
- [powill](#), [Real raised to integer power](#)
- [powll](#), [Power function](#)
- [sinhll](#), [Hyperbolic sine](#)
- [sinll](#), [Circular sine](#)
- [cosll](#), [Circular cosine](#)
- [sqrtll](#), [Square root](#)
- [tanhll](#), [Hyperbolic tangent](#)
- [tanll](#), [Circular tangent](#)
- [cotll](#), [Circular cotangent](#)

```
/*                                     acosh1.c
*
*      Inverse hyperbolic cosine, 128-bit long double precision
*
*
*
*
* SYNOPSIS:
*
* long double x, y, acosh1();
*
* y = acosh1( x );
*
*
*
* DESCRIPTION:
*
* Returns inverse hyperbolic cosine of argument.
```

```

/* If 1 <= x < 1.5, a rational approximation
*
*      sqrt(2z) * P(z)/Q(z)
*
* where z = x-1, is used. Otherwise,
*
* acosh(x) = log( x + sqrt( (x-1)(x+1) ) ).
*
*
*
* ACCURACY:
*
*                                     Relative error:
* arithmetic   domain    # trials     peak       rms
* IEEE         1,3        100,000      4.1e-34     7.3e-35
*
* ERROR MESSAGES:
*
* message          condition      value returned
* acoshl domain    |x| < 1              0.0
*/

/* asinh1.c
*
* Inverse hyperbolic sine, 128-bit long double precision
*
*
* SYNOPSIS:
*
* long double x, y, asinh1();
*
* y = asinh1( x );
*
*
* DESCRIPTION:
*
* Returns inverse hyperbolic sine of argument.
*
* If |x| < 0.5, the function is approximated by a rational
* form x + x**3 P(x)/Q(x). Otherwise,
*
*      asinh(x) = log( x + sqrt(1 + x*x) ).
*
*
* ACCURACY:
*
*                                     Relative error:
* arithmetic   domain    # trials     peak       rms
* IEEE         -2,2        100,000      2.8e-34     6.7e-35
*/

/* asin1.c
*
* Inverse circular sine, 128-bit long double precision
*
*
* SYNOPSIS:
*
* double x, y, asin1();
*
* y = asin1( x );
*
*
* DESCRIPTION:
*
* Returns radian angle between -pi/2 and +pi/2 whose sine is x.
*
* A rational function of the form x + x**3 P(x**2)/Q(x**2)
* is used for |x| in the interval [0, 0.5]. If |x| > 0.5 it is
* transformed by the identity
*
*      asin(x) = pi/2 - 2 asin( sqrt( (1-x)/2 ) ).
*
*
* ACCURACY:
*
*                                     Relative error:
* arithmetic   domain    # trials     peak       rms
* IEEE         -1, 1        100,000      3.7e-34     6.4e-35
*
* ERROR MESSAGES:
*
* message          condition      value returned
* asin domain      |x| > 1              0.0
*/

```

```

/*                                                    acosl()
*
*      Inverse circular cosine, long double precision
*
*
*
* SYNOPSIS:
*
* double x, y, acosl();
*
* y = acosl( x );
*
*
*
* DESCRIPTION:
*
* Returns radian angle between -pi/2 and +pi/2 whose cosine
* is x.
*
* Analytically, acos(x) = pi/2 - asin(x). However if |x| is
* near 1, there is cancellation error in subtracting asin(x)
* from pi/2. Hence if x < -0.5,
*
*     acos(x) = pi - 2.0 * asin( sqrt((1+x)/2) );
*
* or if x > +0.5,
*
*     acos(x) = 2.0 * asin( sqrt((1-x)/2) ).
*
*
* ACCURACY:
*
*                                     Relative error:
* arithmetic   domain   # trials   peak       rms
*   IEEE       -1, 1     100,000    2.1e-34    5.6e-35
*
*
* ERROR MESSAGES:
*
*   message           condition      value returned
* asin domain         |x| > 1         0.0
*/

/*                                                    atanh1.c
*
*      Inverse hyperbolic tangent, 128-bit long double precision
*
*
*
* SYNOPSIS:
*
* long double x, y, atanh1();
*
* y = atanh1( x );
*
*
*
* DESCRIPTION:
*
* Returns inverse hyperbolic tangent of argument in the range
* MINLOGL to MAXLOGL.
*
* If |x| < 0.5, the rational form x + x**3 P(x)/Q(x) is
* employed. Otherwise,
*     atanh(x) = 0.5 * log( (1+x)/(1-x) ).
*
*
*
* ACCURACY:
*
*                                     Relative error:
* arithmetic   domain   # trials   peak       rms
*   IEEE       -1,1     100,000    2.0e-34    4.6e-35
*/

/*                                                    atan1.c
*
*      Inverse circular tangent, 128-bit long double precision
*      (arctangent)
*
*
*
* SYNOPSIS:
*
* long double x, y, atan1();
*
* y = atan1( x );
*
*
*
* DESCRIPTION:
*
* Returns radian angle between -pi/2 and +pi/2 whose tangent

```



```

*
* y = coshl( x );
*
*
* DESCRIPTION:
*
* Returns hyperbolic cosine of argument in the range MINLOGL to
* MAXLOGL.
*
* cosh(x) = ( exp(x) + exp(-x) )/2.
*
*
* ACCURACY:
*
*                Relative error:
* arithmetic   domain   # trials   peak       rms
*   IEEE      +-10000    26,000     2.5e-34    8.6e-35
*
* ERROR MESSAGES:
*
*   message      condition      value returned
* cosh overflow  |x| > MAXLOGL    MAXNUML
*
*/

```

```

/*                                     exp10l.c
*
*   Base 10 exponential function, long double precision
*   (Common antilogarithm)
*
*
* SYNOPSIS:
*
* long double x, y, exp10l()
*
* y = exp10l( x );
*
*
* DESCRIPTION:
*
* Returns 10 raised to the x power.
*
* Range reduction is accomplished by expressing the argument
* as 10**x = 2**n 10**f, with |f| < 0.5 log10(2).
* The Pade' form
*
*   1 + 2x P(x**2)/( Q(x**2) - P(x**2) )
*
* is used to approximate 10**f.
*
*
* ACCURACY:
*
*                Relative error:
* arithmetic   domain   # trials   peak       rms
*   IEEE      +-4900    100,000     2.1e-34    4.7e-35
*
* ERROR MESSAGES:
*
*   message      condition      value returned
* exp10l underflow  x < -MAXL10    0.0
* exp10l overflow   x > MAXL10    MAXNUM
*
* IEEE arithmetic: MAXL10 = 4932.0754489586679023819
*
*/

```

```

/*                                     exp2l.c
*
*   Base 2 exponential function, 128-bit long double precision
*
*
* SYNOPSIS:
*
* long double x, y, exp2l();
*
* y = exp2l( x );
*
*
* DESCRIPTION:
*
* Returns 2 raised to the x power.
*
* Range reduction is accomplished by separating the argument
* into an integer k and fraction f such that
*
*   x    k  f
* 2  = 2  2.

```

```

*
* A Pade' form
*
*  $1 + 2x P(x^{**2}) / (Q(x^{**2}) - x P(x^{**2}))$ 
*
* approximates  $2^{**x}$  in the basic range  $[-0.5, 0.5]$ .
*
*
* ACCURACY:
*
*                Relative error:
* arithmetic  domain  # trials   peak       rms
*   IEEE      +-16300   100,000    2.0e-34    4.8e-35
*
* See exp.c for comments on error amplification.
*
*
* ERROR MESSAGES:
*
* message      condition      value returned
* exp2l underflow  x < -16382         0.0
* exp2l overflow   x >= 16384        MAXNUM
*
*/

/*                                     expl.c
*
*      Exponential function, 128-bit long double precision
*
*
* SYNOPSIS:
*
* long double x, y, expl();
*
* y = expl( x );
*
*
* DESCRIPTION:
*
* Returns e (2.71828...) raised to the x power.
*
* Range reduction is accomplished by separating the argument
* into an integer k and fraction f such that
*
*      x    k  f
*      e  = 2  e.
*
* A Pade' form of degree 2/3 is used to approximate  $\exp(f) - 1$ 
* in the basic range  $[-0.5 \ln 2, 0.5 \ln 2]$ .
*
*
* ACCURACY:
*
*                Relative error:
* arithmetic  domain  # trials   peak       rms
*   IEEE      +-MAXLOG  100,000    2.6e-34    8.6e-35
*
* Error amplification in the exponential function can be
* a serious matter. The error propagation involves
*  $\exp(X(1+\delta)) = \exp(X) (1 + X\delta + \dots)$ ,
* which shows that a 1 lsb error in representing X produces
* a relative error of X times 1 lsb in the function.
* While the routine gives an accurate result for arguments
* that are exactly represented by a long double precision
* computer number, the result contains amplified roundoff
* error for large arguments not exactly represented.
*
*
* ERROR MESSAGES:
*
* message      condition      value returned
* exp underflow  x < MINLOG         0.0
* exp overflow   x > MAXLOG        MAXNUM
*
*/

/*                                     expm1l.c
*
*      Exponential function, minus 1
*      128-bit long double precision
*
*
* SYNOPSIS:
*
* long double x, y, expm1l();
*
* y = expm1l( x );
*
*
* DESCRIPTION:

```

```

*
* Returns e (2.71828...) raised to the x power, minus 1.
*
* Range reduction is accomplished by separating the argument
* into an integer k and fraction f such that
*
*      x   k   f
*      e = 2  e.
*
* An expansion  $x + .5 x^2 + x^3 R(x)$  approximates  $\exp(f) - 1$ 
* in the basic range  $[-0.5 \ln 2, 0.5 \ln 2]$ .
*
*
* ACCURACY:
*
*                Relative error:
* arithmetic  domain  # trials  peak       rms
*   IEEE      -79,+MAXLOG   100,000   1.7e-34    4.5e-35
*
* ERROR MESSAGES:
*
*   message      condition      value returned
* expm1 overflow  x > MAXLOG        MAXNUM
*
*/

/*                                ceil()
*                                floorl()
*                                frexpl()
*                                ldexpl()
*                                fabsl()
*                                signbitl()
*                                isnanl()
*                                isfinitel()
*
* Floating point numeric utilities
*
*
* SYNOPSIS:
*
* long double x, y;
* long double ceil(), floorl(), frexpl(), ldexpl(), fabsl();
* int signbitl(), isnanl(), isfinitel();
* int expnt, n;
*
* y = floorl(x);
* y = ceil(x);
* y = frexpl( x, &expnt );
* y = ldexpl( x, n );
* y = fabsl( x );
*
*
* DESCRIPTION:
*
* All four routines return a long double precision floating point
* result.
*
* floorl() returns the largest integer less than or equal to x.
* It truncates toward minus infinity.
*
* ceil() returns the smallest integer greater than or equal
* to x. It truncates toward plus infinity.
*
* frexpl() extracts the exponent from x. It returns an integer
* power of two to expnt and the significand between 0.5 and 1
* to y. Thus  $x = y * 2^{expn}$ .
*
* ldexpl() multiplies x by  $2^n$ .
*
* fabsl() returns the absolute value of its argument.
*
* signbitl(x) returns 1 if the sign bit of x is 1, else 0.
*
* These functions are part of the standard C run time library
* for some but not all C compilers. The ones supplied are
* written in C for IEEE arithmetic. They should
* be used only if your compiler library does not already have
* them.
*
* The IEEE versions assume that denormal numbers are implemented
* in the arithmetic. Some modifications will be required if
* the arithmetic has abrupt rather than gradual underflow.
*/

/*                                ieee.c
*
* Extended precision IEEE binary floating point arithmetic routines
*
* Numbers are stored in C language as arrays of 16-bit unsigned
* short integers. The arguments of the routines are pointers to
* the arrays.
*
*
* External e type data structure, simulates Intel 8087 chip

```

```

* temporary real format but possibly with a larger significand:
*
*     NE-1 significand words  (least significant word first,
*                             most significant bit is normally set)
*     exponent                (value = EXONE for 1.0,
*                             top bit is the sign)
*
*
* Internal data structure of a number (a "word" is 16 bits):
*
* ei[0]      sign word      (0 for positive, 0xffff for negative)
* ei[1]      biased exponent (value = EXONE for the number 1.0)
* ei[2]      high guard word (always zero after normalization)
* ei[3]
* to ei[NI-2] significand   (NI-4 significand words,
*                             most significant word first,
*                             most significant bit is set)
* ei[NI-1]   low guard word (0x8000 bit is rounding place)
*
*
*
*           Routines for external format numbers
*
* asctoe( string, e )      ASCII string to extended double e type
* asctoe64( string, &d )  ASCII string to long double
* asctoe53( string, &d )  ASCII string to double
* asctoe24( string, &f )  ASCII string to single
* asctoeg( string, e, prec ) ASCII string to specified precision
* e24toe( &f, e )         IEEE single precision to e type
* e53toe( &d, e )         IEEE double precision to e type
* e64toe( &d, e )         IEEE long double precision to e type
* eabs(e)                 absolute value
* eadd( a, b, c )          c = b + a
* eclear(e)                e = 0
* ecmp( a, b )             Returns 1 if a > b, 0 if a == b,
*                           -1 if a < b, -2 if either a or b is a NaN.
* ediv( a, b, c )          c = b / a
* efloor( a, b )           truncate to integer, toward -infinity
* efrac( a, exp, s )       extract exponent and significand
* eifrac( e, &l, frac )    e to long integer and e type fraction
* eifrac( e, &l, frac )    e to unsigned long integer and e type fraction
* einfin( e )              set e to infinity, leaving its sign alone
* eldexp( a, n, b )        multiply by 2**n
* emov( a, b )             b = a
* emul( a, b, c )          c = b * a
* eneg(e)                  e = -e
* eround( a, b )           b = nearest integer value to a
* esub( a, b, c )          c = b - a
* e24toasc( &f, str, n )  single to ASCII string, n digits after decimal
* e53toasc( &d, str, n )  double to ASCII string, n digits after decimal
* e64toasc( &d, str, n )  long double to ASCII string
* etoasc( e, str, n )      e to ASCII string, n digits after decimal
* etoe24( e, &f )          convert e type to IEEE single precision
* etoe53( e, &d )          convert e type to IEEE double precision
* etoe64( e, &d )          convert e type to IEEE long double precision
* ltoe( &l, e )            long (32 bit) integer to e type
* ultoe( &l, e )           unsigned long (32 bit) integer to e type
* eisneg( e )              1 if sign bit of e != 0, else 0
* eisinf( e )              1 if e has maximum exponent (non-IEEE)
*                           or is infinite (IEEE)
* eisnan( e )              1 if e is a NaN
* esqrt( a, b )            b = square root of a
*
*
*           Routines for internal format numbers
*
* eaddm( ai, bi )          add significands, bi = bi + ai
* ecleaz(ei)               ei = 0
* ecleazs(ei)              set ei = 0 but leave its sign alone
* ecmpm( ai, bi )          compare significands, return 1, 0, or -1
* edivm( ai, bi )          divide significands, bi = bi / ai
* emdnorm(ai,l,s,exp)      normalize and round off
* emovi( a, ai )           convert external a to internal ai
* emovo( ai, a )           convert internal ai to external a
* emovz( ai, bi )          bi = ai, low guard word of bi = 0
* emulm( ai, bi )          multiply significands, bi = bi * ai
* enormlz(ei)              left-justify the significand
* eshdn1( ai )             shift significand and guards down 1 bit
* eshdn8( ai )             shift down 8 bits
* eshdn6( ai )             shift down 16 bits
* eshift( ai, n )          shift ai n bits up (or down if n < 0)
* eshup1( ai )             shift significand and guards up 1 bit
* eshup8( ai )             shift up 8 bits
* eshup6( ai )             shift up 16 bits
* esubm( ai, bi )          subtract significands, bi = bi - ai
*
*
* The result is always normalized and rounded to NI-4 word precision
* after each arithmetic operation.
*
* Exception flags are NOT fully supported.
*
* Define INFINITIES in mconf.h for support of infinity; otherwise a
* saturation arithmetic is implemented.
*
* Define NANS for support of Not-a-Number items; otherwise the
* arithmetic will never produce a NaN output, and might be confused
* by a NaN input.
* If NaN's are supported, the output of ecmp(a,b) is -2 if
* either a or b is a NaN. This means asking if(ecmp(a,b) < 0)

```



```

/*                                                    j0l.c
*
*      Bessel function of order zero
*
*
*
* SYNOPSIS:
*
* long double x, y, j0l();
*
* y = j0l( x );
*
*
* DESCRIPTION:
*
* Returns Bessel function of first kind, order zero of the argument.
*
* The domain is divided into two major intervals [0, 2] and
* (2, infinity). In the first interval the rational approximation
* is  $J_0(x) = 1 - x^2 / 4 + x^4 R(x^2)$ 
* The second interval is further partitioned into eight equal segments
* of  $1/x$ .
*
*  $J_0(x) = \sqrt{2/(pi\ x)} (P_0(x)\ cos(X) - Q_0(x)\ sin(X)),$ 
*  $X = x - pi/4,$ 
*
* and the auxiliary functions are given by
*
*  $J_0(x)\cos(X) + Y_0(x)\sin(X) = \sqrt{2/(pi\ x)} P_0(x),$ 
*  $P_0(x) = 1 + 1/x^2 R(1/x^2)$ 
*
*  $Y_0(x)\cos(X) - J_0(x)\sin(X) = \sqrt{2/(pi\ x)} Q_0(x),$ 
*  $Q_0(x) = 1/x (-.125 + 1/x^2 R(1/x^2))$ 
*
*
*
* ACCURACY:
*
*
* Absolute error:
*


| arithmetic | domain | # trials | peak    | rms     |
|------------|--------|----------|---------|---------|
| IEEE       | 0, 30  | 100000   | 1.7e-34 | 2.4e-35 |


*/

```

```
/*                                                    j111.c
 *
 *      Bessel function of order one
 *
 *
 *
 *
 * SYNOPSIS:
 *
 * long double x, y, j11();
```

```

*
* y = j1l( x );
*
*
* DESCRIPTION:
*
* Returns Bessel function of first kind, order one of the argument.
*
* The domain is divided into two major intervals [0, 2] and
* (2, infinity). In the first interval the rational approximation is
*  $J_1(x) = .5x + x \cdot x^2 R(x^2)$ 
*
* The second interval is further partitioned into eight equal segments
* of  $1/x$ .
*  $J_1(x) = \sqrt{2/(\pi x)} (P_1(x) \cos(X) - Q_1(x) \sin(X))$ ,
*  $X = x - 3 \pi / 4$ ,
*
* and the auxiliary functions are given by
*
*  $J_1(x)\cos(X) + Y_1(x)\sin(X) = \sqrt{2/(\pi x)} P_1(x)$ ,
*  $P_1(x) = 1 + 1/x^2 R(1/x^2)$ 
*
*  $Y_1(x)\cos(X) - J_1(x)\sin(X) = \sqrt{2/(\pi x)} Q_1(x)$ ,
*  $Q_1(x) = 1/x (.375 + 1/x^2 R(1/x^2))$ .
*
*
*
* ACCURACY:
*
*
* Absolute error:
* arithmetic domain # trials peak rms
* IEEE 0, 30 100000 2.8e-34 2.7e-35
*
*
*/

```

```

/*
/*
/*
/* Bessel function of the second kind, order one
/*
/*
/* SYNOPSIS:
/*
/* double x, y, y1l();
/*
/* y = y1l( x );
/*
/*
/* DESCRIPTION:
/*
/* Returns Bessel function of the second kind, of order
/* one, of the argument.
/*
/* The domain is divided into two major intervals [0, 2] and
/* (2, infinity). In the first interval the rational approximation is
/*  $Y_1(x) = 2/\pi * (\log(x) * J_1(x) - 1/x) + x R(x^2)$  .
/* In the second interval the approximation is the same as for  $J_1(x)$ , and
/*  $Y_1(x) = \sqrt{2/(\pi x)} (P_1(x) \sin(X) + Q_1(x) \cos(X))$ ,
/*  $X = x - 3 \pi / 4$ .
/*
/* ACCURACY:
/*
/* Absolute error, when  $y_0(x) < 1$ ; else relative error:
/*
/* arithmetic domain # trials peak rms
/* IEEE 0, 30 100000 2.7e-34 2.9e-35
/*
*/

```

```

/*
/*
/*
/* Bessel function of integer order
/*
/*
/* SYNOPSIS:
/*
/* int n;
/* long double x, y, jnl();
/*
/* y = jnl( n, x );
/*
/*
/* DESCRIPTION:
/*
/* Returns Bessel function of order n, where n is a
/* (possibly negative) integer.
/*
/* The ratio of  $j_n(x)$  to  $j_0(x)$  is computed by backward
/* recurrence. First the ratio  $j_n/j_{n-1}$  is found by a
/* continued fraction expansion. Then the recurrence
/* relating successive orders is applied until  $j_0$  or  $j_1$  is

```

```

* reached.
*
* If n = 0 or 1 the routine for j0 or j1 is called
* directly.
*
*
*
* ACCURACY:
*
*
* Absolute error:
* arithmetic domain # trials peak rms
* IEEE -30, 30 10000 2.6e-34 4.6e-35
*
*
* Not suitable for large n or x.
*
*/

```

```

/* lgammall.c
*
* Natural logarithm of gamma function
*
*
* SYNOPSIS:
*
* long double x, y, lgammal();
* extern int sgngam;
*
* y = lgammal(x);
*
*
* DESCRIPTION:
*
* Returns the base e (2.718...) logarithm of the absolute
* value of the gamma function of the argument.
* The sign (+1 or -1) of the gamma function is returned in a
* global (extern) variable named sgngam.
*
* The positive domain is partitioned into numerous segments for approximation.
* For x > 10,
*   log gamma(x) = (x - 0.5) log(x) - x + log sqrt(2 pi) + 1/x R(1/x^2)
* Near the minimum at x = x0 = 1.46... the approximation is
*   log gamma(x0 + z) = log gamma(x0) + z^2 P(z)/Q(z)
* for small z.
* Elsewhere between 0 and 10,
*   log gamma(n + z) = log gamma(n) + z P(z)/Q(z)
* for various selected n and small z.
*
* The cosecant reflection formula is employed for negative arguments.
*
* Arguments greater than MAXLGML (10^4928) return MAXNUML.
*
*
* ACCURACY:
*
*
* arithmetic domain # trials peak rms
* Relative error:
* IEEE 10, 30 100000 3.9e-34 9.8e-35
* IEEE 0, 10 100000 3.8e-34 5.3e-35
* Absolute error:
* IEEE -10, 0 100000 8.0e-34 8.0e-35
* IEEE -30, -10 100000 4.4e-34 1.0e-34
* IEEE -100, 100 100000 1.0e-34
*
* The absolute error criterion is the same as relative error
* when the function magnitude is greater than one but it is absolute
* when the magnitude is less than one.
*
*/

```

```

/* log10l.c
*
* Common logarithm, long double precision
*
*
* SYNOPSIS:
*
* long double x, y, log10l();
*
* y = log10l( x );
*
*
* DESCRIPTION:
*
* Returns the base 10 logarithm of x.
*
* The argument is separated into its exponent and fractional
* parts. If the exponent is between -1 and +1, the logarithm
* of the fraction is approximated by
*
*   log(1+x) = x - 0.5 x**2 + x**3 P(x)/Q(x).

```

```

/*
 * Otherwise, setting  $z = 2(x-1)/(x+1)$ ,
 *
 *  $\log(x) = z + z^3 P(z)/Q(z).$ 
 *
 *
 *
 * ACCURACY:
 *
 *
 * Relative error:
 * arithmetic domain # trials peak rms
 * IEEE 0.5, 2.0 30000 2.3e-34 4.9e-35
 * IEEE exp(+10000) 30000 1.0e-34 4.1e-35
 *
 * In the tests over the interval exp(+10000), the logarithms
 * of the random arguments were uniformly distributed over
 * [-10000, +10000].
 *
 * ERROR MESSAGES:
 *
 * log singularity: x = 0; returns MINLOG
 * log domain: x < 0; returns MINLOG
 */

/* log1pl.c
 *
 * Relative error logarithm
 * Natural logarithm of 1+x, 128-bit long double precision
 *
 *
 * SYNOPSIS:
 *
 * long double x, y, log1pl();
 *
 * y = log1pl( x );
 *
 * DESCRIPTION:
 *
 * Returns the base e (2.718...) logarithm of 1+x.
 *
 * The argument 1+x is separated into its exponent and fractional
 * parts. If the exponent is between -1 and +1, the logarithm
 * of the fraction is approximated by
 *
 *  $\log(1+x) = x - 0.5 x^2 + x^3 P(x)/Q(x).$ 
 *
 * Otherwise, setting  $z = 2(x-1)/(x+1)$ ,
 *
 *  $\log(x) = z + z^3 P(z)/Q(z).$ 
 *
 *
 *
 * ACCURACY:
 *
 *
 * Relative error:
 * arithmetic domain # trials peak rms
 * IEEE -1, 8 100000 1.9e-34 4.3e-35
 */

/* log2l.c
 *
 * Base 2 logarithm, long double precision
 *
 *
 * SYNOPSIS:
 *
 * long double x, y, log2l();
 *
 * y = log2l( x );
 *
 * DESCRIPTION:
 *
 * Returns the base 2 logarithm of x.
 *
 * The argument is separated into its exponent and fractional
 * parts. If the exponent is between -1 and +1, the (natural)
 * logarithm of the fraction is approximated by
 *
 *  $\log(1+x) = x - 0.5 x^{**2} + x^{**3} P(x)/Q(x).$ 
 *
 * Otherwise, setting  $z = 2(x-1)/(x+1)$ ,
 *
 *  $\log(x) = z + z^{**3} P(z)/Q(z).$ 
 *
 *
 *
 * ACCURACY:
 *
 *
 * Relative error:
 * arithmetic domain # trials peak rms

```

```

*      IEEE      0.5, 2.0      100,000      1.3e-34      4.5e-35
*      IEEE      exp(+10000)  100,000      9.6e-35      4.0e-35
*
* In the tests over the interval exp(+10000), the logarithms
* of the random arguments were uniformly distributed over
* [-10000, +10000].
*
* ERROR MESSAGES:
*
* log singularity:  x = 0; returns MINLOG
* log domain:      x < 0; returns MINLOG
*/

```

```

/*                                          logl.c
*
*      Natural logarithm, long double precision
*
*
* SYNOPSIS:
*
* long double x, y, logl();
*
* y = logl( x );
*
*
* DESCRIPTION:
*
* Returns the base e (2.718...) logarithm of x.
*
* The argument is separated into its exponent and fractional
* parts.  If the exponent is between -1 and +1, the logarithm
* of the fraction is approximated by
*
*      log(1+x) = x - 0.5 x**2 + x**3 P(x)/Q(x).
*
* Otherwise, setting  z = 2(x-1)/(x+1),
*
*      log(x) = z + z**3 P(z)/Q(z).
*
*
* ACCURACY:
*
*                                     Relative error:
* arithmetic  domain  # trials  peak       rms
*   IEEE      exp(+MAXLOGL) 36,000    9.5e-35    4.1e-35
*
* ERROR MESSAGES:
*
* log singularity:  x = 0; returns MINLOGL
* log domain:      x < 0; returns MINLOGL
*/

```

```

/*                                          ndtr1l.c
*
*      Normal distribution function
*      128-bit long double version
*
*
* SYNOPSIS:
*
* long double x, y, ndtr1l();
*
* y = ndtr1l( x );
*
*
* DESCRIPTION:
*
* Returns the area under the Gaussian probability density
* function, integrated from minus infinity to x:
*
*
*
*
*

$$\text{ndtr}(x) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^x \exp\left(-t^2/2\right) dt$$

*
*

$$= (1 + \text{erf}(z)) / 2$$


$$= \text{erfc}(z) / 2$$

*
* where z = x/sqrt(2). Computation is via the functions
* erf and erfc with care to avoid error amplification in computing exp(-x^2).
*
*
* ACCURACY:
*
*                                     Relative error:
* arithmetic  domain  # trials  peak       rms
*   IEEE      -13,0    50000    7.7e-34    1.7e-34

```

```
*
*
* ERROR MESSAGES:
*
*   message          condition          value returned
* erfcl underflow   x^2 / 2 > MAXLOGL      0.0
*
*/
```

```
*
*      Error function
```

```
*
* long double x, y, erfl();
```

*
*

```
*
* The integral is
```

```
*
* The magnitude of x is limited to about 106.56 for IEEE
* arithmetic; 1 or -1 is returned outside this range.
*
* For  $0 \leq |x| < 1$ , erf(x) is computed by rational approximations; otherwise
* erf(x) = 1 - erfc(x).
```

```
*
*                               Relative error:
* arithmetic      domain      # trials      peak      rms
*      IEEE      0,1          50000      1.5e-34      4.4e-35
*
*/
```

```
*
* Complementary error function
```

```
*
* long double x, y, erfc1();
```

*
*

*
*

```
*
*
* For small x, erfc(x) = 1 - erf(x); otherwise rational
* approximations are computed.
```

```

*
*
*          Relative error:
* arithmetic    domain    # trials    peak      rms
*      IEEE      0,13      100000    5.8e-34    1.5e-34
*      IEEE      6,106.56  100000    5.9e-34    1.5e-34
*
*

```

```

* ERROR MESSAGES:
*
*   message          condition          value returned
* erfcl underflow    x^2 > MAXLOGL      0.0
*
*/

```

```

/*                                          mtherr.c
*
*   Library common error handling routine
*
*
* SYNOPSIS:
*
* char *fctnam;
* int code;
* void mtherr();
*
* mtherr( fctnam, code );
*
*
* DESCRIPTION:
*
* This routine may be called to report one of the following
* error conditions (in the include file mconf.h).
*
*   Mnemonic          Value          Significance
*
*   DOMAIN            1            argument domain error
*   SING              2            function singularity
*   OVERFLOW          3            overflow range error
*   UNDERFLOW        4            underflow range error
*   TLOSS             5            total loss of precision
*   PLOSS             6            partial loss of precision
*   EDOM              33           Unix domain error code
*   ERANGE            34           Unix range error code
*
* The default version of the file prints the function name,
* passed to it by the pointer fctnam, followed by the
* error condition. The display is directed to the standard
* output device. The routine then returns to the calling
* program. Users may wish to modify the program to abort by
* calling exit() under severe error conditions such as domain
* errors.
*
* Since all error conditions pass control to this function,
* the display may be easily changed, eliminated, or directed
* to an error logging device.
*
* SEE ALSO:
*
* mconf.h
*
*/

```

```

/*                                          polevll.c
*                                          plevll.c
*
*   Evaluate polynomial
*
*
* SYNOPSIS:
*
* int N;
* long double x, y, coef[N+1], polevl[];
*
* y = polevll( x, coef, N );
*
*
* DESCRIPTION:
*
* Evaluates polynomial of degree N:
*
*      2      N
* y = C0 + C1 x + C2 x + ... + CN x
*
* Coefficients are stored in reverse order:
*
* coef[0] = CN , ..., coef[N] = C0 .
*
* The function plevll() assumes that coef[N] = 1.0 and is
* omitted from the array. Its calling arguments are
* otherwise the same as polevll().
*
*
* SPEED:
*
* In the interest of speed, there are no checks for out
* of bounds arithmetic. This routine is used by most of

```

```

* the functions in the library. Depending on available
* equipment features, the user may wish to rewrite the
* program in microcode or assembly language.

```

```

*/

```

```

/*                                     powil.c

```

```

*      Real raised to integer power, long double precision

```

```

* SYNOPSIS:

```

```

* long double x, y, powil();
* int n;

```

```

* y = powil( x, n );

```

```

* DESCRIPTION:

```

```

* Returns argument x raised to the nth power.
* The routine efficiently decomposes n as a sum of powers of
* two. The desired power is a product of two-to-the-kth
* powers of x. Thus to compute the 32767 power of x requires
* 28 multiplications instead of 32767 multiplications.

```

```

* ACCURACY:

```

```

*                                     Relative error:
* arithmetic   x domain   n domain # trials   peak       rms
*   IEEE       .001,1000  -1022,1023  100,000   7.5e-32    1.4e-32
*   IEEE       .99,1.01   0,8700    100,000   4.6e-31    9.1e-32

```

```

* Returns MAXNUM on overflow, zero on underflow.

```

```

*/

```

```

/*                                     powl.c

```

```

*      Power function, long double precision

```

```

* SYNOPSIS:

```

```

* long double x, y, z, powl();

```

```

* z = powl( x, y );

```

```

* DESCRIPTION:

```

```

* Computes x raised to the yth power. For noninteger y,

```

```

*      x^y = exp2( y log2(x) ).

```

```

* using the base 2 logarithm and exponential functions. If y
* is an integer, |y| < 32768, the function is computed by powil.

```

```

* ACCURACY:

```

```

* The relative error of pow(x,y) can be estimated
* by y dl ln(2), where dl is the absolute error of
* the internally computed base 2 logarithm.

```

```

*                                     Relative error:
* arithmetic   domain     # trials   peak       rms
*   IEEE       +-1000     100,000    1.0e-30    1.4e-31
* .001 < x < 1000, with log(x) uniformly distributed.
* -1000 < y < 1000, y uniformly distributed.

```

```

*   IEEE       0,8700     100,000    1.4e-30    3.1e-31
* 0.99 < x < 1.01, 0 < y < 8700, uniformly distributed.

```

```

* ERROR MESSAGES:

```

```

* message      condition      value returned
* pow overflow  x^y > MAXNUM             MAXNUM
* pow underflow x^y < 1/MAXNUM      0.0
* pow domain    x<0 and y noninteger 0.0

```

```

*/

```



```

/*                                     sinhl.c
*
*      Hyperbolic sine, 128-bit long double precision
*
*
* SYNOPSIS:
*
* long double x, y, sinhl();
* y = sinhl( x );
*
*
* DESCRIPTION:
*
* Returns hyperbolic sine of argument in the range MINLOGL to
* MAXLOGL.
*
* The range is partitioned into two segments.  If  $|x| \leq 1$ , a
* rational function of the form  $x + x^3 P(x)/Q(x)$  is employed.
* Otherwise the calculation is  $\sinh(x) = (\exp(x) - \exp(-x))/2$ .
*
*
* ACCURACY:
*
*                                     Relative error:
* arithmetic   domain    # trials   peak       rms
* IEEE         -2,2      100,000    4.1e-34    7.9e-35
*/

```

```

/*                                     sinl.c
*
*      Circular sine, long double precision
*
*
* SYNOPSIS:
*
* long double x, y, sinl();
* y = sinl( x );
*
*
* DESCRIPTION:
*
* Range reduction is into intervals of  $\pi/4$ .  The reduction
* error is nearly eliminated by contriving an extended precision
* modular arithmetic.
*
* Two polynomial approximating functions are employed.
* Between 0 and  $\pi/4$  the sine is approximated by the Cody
* and Waite polynomial form
*  $x + x^3 P(x^2)$  .
* Between  $\pi/4$  and  $\pi/2$  the cosine is represented as
*  $1 - .5 x^2 + x^4 Q(x^2)$  .
*
*
* ACCURACY:
*
*                                     Relative error:
* arithmetic   domain    # trials   peak       rms
* IEEE         +-3.6e16   100,000    2.0e-34    5.3e-35
*
* ERROR MESSAGES:
*
* message      condition      value returned
* sin total loss   $x > 2^{55}$          0.0
*/

```

```

/*                                     cosl.c
*
*      Circular cosine, long double precision
*
*
* SYNOPSIS:
*
* long double x, y, cosl();
* y = cosl( x );
*
*
* DESCRIPTION:
*
* Range reduction is into intervals of  $\pi/4$ .  The reduction
* error is nearly eliminated by contriving an extended precision
* modular arithmetic.
*
* Two polynomial approximating functions are employed.
* Between 0 and  $\pi/4$  the cosine is approximated by

```

```

*      1 - .5 x^2 + x^4 Q(x^2) .
* Between pi/4 and pi/2 the sine is represented by the Cody
* and Waite polynomial form
*      x + x^3 P(x^2) .
*
*
* ACCURACY:
*
*                Relative error:
* arithmetic  domain  # trials   peak      rms
*   IEEE      +-3.6e16   100,000    2.0e-34    5.2e-35
*
* ERROR MESSAGES:
*
*   message          condition      value returned
* cos total loss     x > 2^55         0.0
*/

```

```

/*                                     sqrtl.c
*
*      Square root, long double precision
*
*
* SYNOPSIS:
*
* long double x, y, sqrtl();
* y = sqrtl( x );
*
*
* DESCRIPTION:
*
* Returns the square root of x.
*
* Range reduction involves isolating the power of two of the
* argument and using a polynomial approximation to obtain
* a rough value for the square root. Then Heron's iteration
* is used three times to converge to an accurate value.
*
* Note, some arithmetic coprocessors such as the 8087 and
* 68881 produce correctly rounded square roots, which this
* routine will not.
*
* ACCURACY:
*
*                Relative error:
* arithmetic  domain  # trials   peak      rms
*   IEEE      0,10    30000     8.1e-20    3.1e-20
*
* ERROR MESSAGES:
*
*   message          condition      value returned
* sqrt domain        x < 0          0.0
*/

```

```

/*                                     tanhl.c
*
*      Hyperbolic tangent, 128-bit long double precision
*
*
* SYNOPSIS:
*
* long double x, y, tanhl();
* y = tanhl( x );
*
*
* DESCRIPTION:
*
* Returns hyperbolic tangent of argument in the range MINLOGL to
* MAXLOGL.
*
* A rational function is used for |x| < 0.625. The form
* x + x**3 P(x)/Q(x) of Cody & Waite is employed.
* Otherwise,
*      tanh(x) = sinh(x)/cosh(x) = 1 - 2/(exp(2x) + 1).
*
*
* ACCURACY:
*
*                Relative error:
* arithmetic  domain  # trials   peak      rms
*   IEEE      -2,2    100,000    2.1e-34    4.5e-35
*/

```

```

/*                                                    tan1.c
*
*      Circular tangent, 128-bit long double precision
*
*
* SYNOPSIS:
*
* long double x, y, tanl();
* y = tanl( x );
*
*
* DESCRIPTION:
*
* Returns the circular tangent of the radian argument x.
*
* Range reduction is modulo pi/4. A rational function
*  $x + x^3 P(x^2)/Q(x^2)$ 
* is employed in the basic interval [0, pi/4].
*
*
* ACCURACY:
*
*
*          Relative error:
* arithmetic domain # trials peak rms
* IEEE +-3.6e16 100,000 3.0e-34 7.2e-35
*
* ERROR MESSAGES:
*
* message condition value returned
* tan total loss x > 2^55 0.0
*/

```

```

/*                                                    cot1.c
*
*      Circular cotangent, long double precision
*
*
* SYNOPSIS:
*
* long double x, y, cotl();
* y = cotl( x );
*
*
* DESCRIPTION:
*
* Returns the circular cotangent of the radian argument x.
*
* Range reduction is modulo pi/4. A rational function
*  $x + x^3 P(x^2)/Q(x^2)$ 
* is employed in the basic interval [0, pi/4].
*
*
* ACCURACY:
*
*
*          Relative error:
* arithmetic domain # trials peak rms
* IEEE +-3.6e16 100,000 2.9e-34 7.2e-35
*
* ERROR MESSAGES:
*
* message condition value returned
* cot total loss x > 2^55 0.0
* cot singularity x = 0 MAXNUM
*/

```

[To Cephes home page www.moshier.net](http://www.moshier.net):

Last update: 27 January 2002